



Interprocess Communication





IP Fragmentation

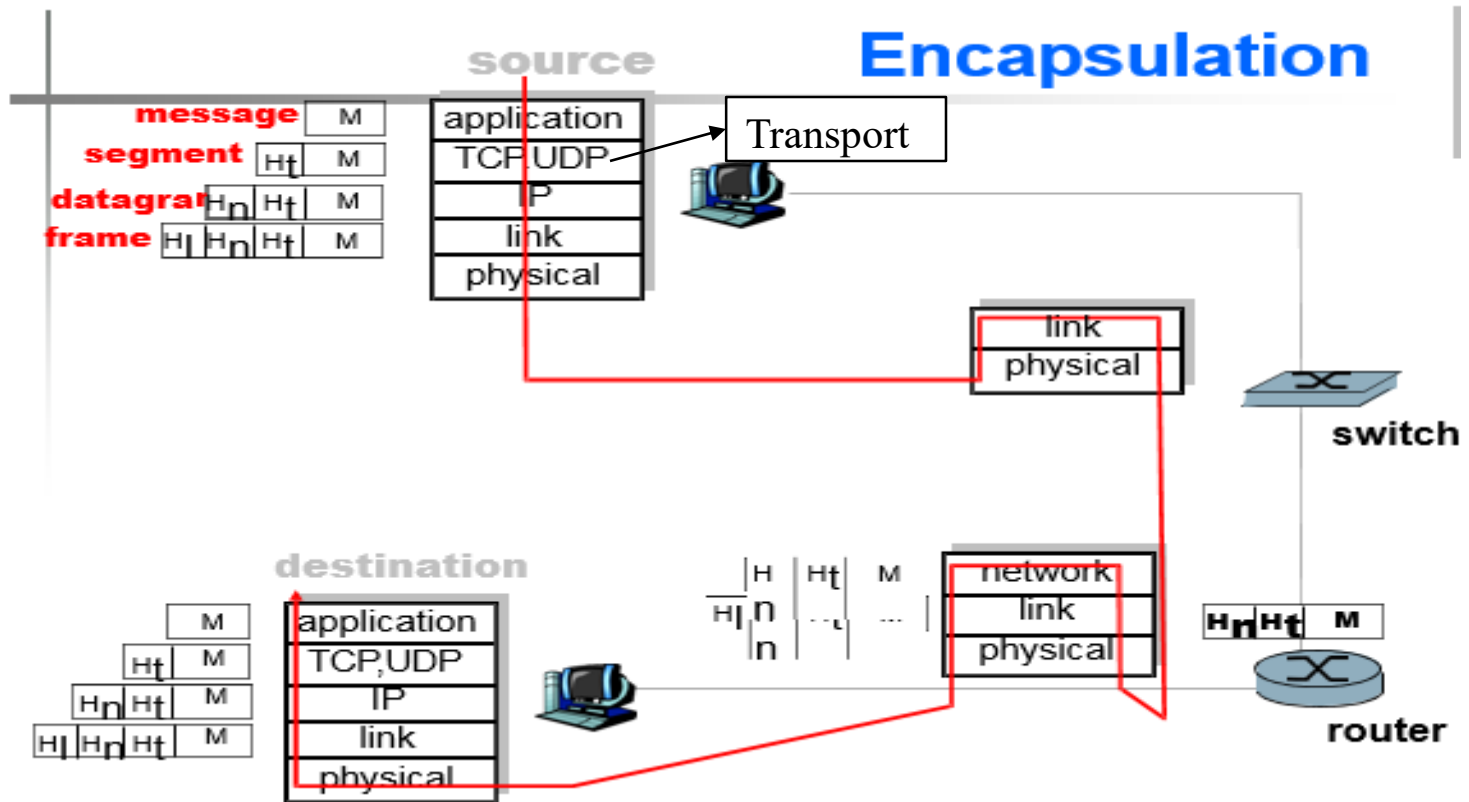
TCP/IP Layers

TCP/IP Protocols

Application Layer	HTTP	FTP	Telnet	SMTP	DNS
Transport Layer	TCP			UDP	
Network Layer	IP		ARP	ICMP	IGMP
Network Interface Layer	Ethernet	Token Ring		Other Link-Layer Protocols	
Physical	Bits of data				

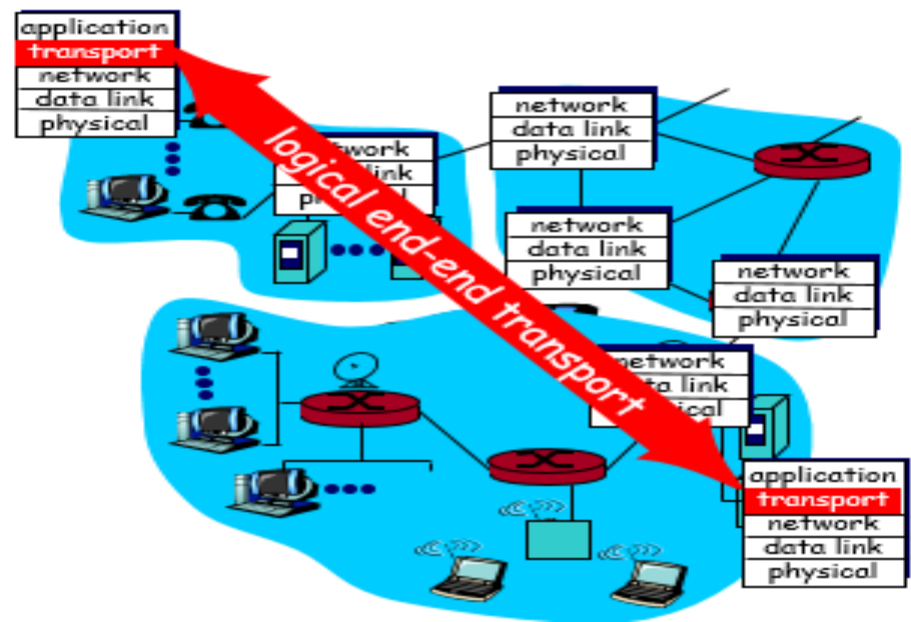
TCP/IP Model

■ TCP/IP Model



Transport Protocols

- Provide *logical communication* between application processes running on different hosts
- Run on end hosts
 - Sender: breaks application messages into **segments**, and passes to network layer
 - Receiver: reassembles segments into messages, passes to application layer
- Multiple transport protocol available to applications
 - Internet: TCP and UDP





Internet Transport Protocols

- User Datagram Protocol (UDP)
 - Provides: Unreliable, unordered delivery of segments.
- Transmission Control Protocol (TCP)
 - Provides: Reliable, in-order delivery of segments.
 - TCP includes:
 - Connection set-up (3-way handshake)
 - Flow Control
 - Congestion Control (packets lost)
 - Discarding of corrupted packets
 - Retransmission of lost packets
- **NOTE:** Neither protocol provides:
 - Delay guarantees
 - Bandwidth guarantees



The characteristics of interprocess communication

- **Communication**: One process sends a message to a destination and another process at the destination receives the message.
- Communication between the sending and receiving processes may be either **synchronous** (blocking) or **asynchronous** (non-blocking).
- **Queue**: A queue is associated with each message destination.
 - Sending processes cause messages to be added to remote queues and receiving processes remove messages from local queues.

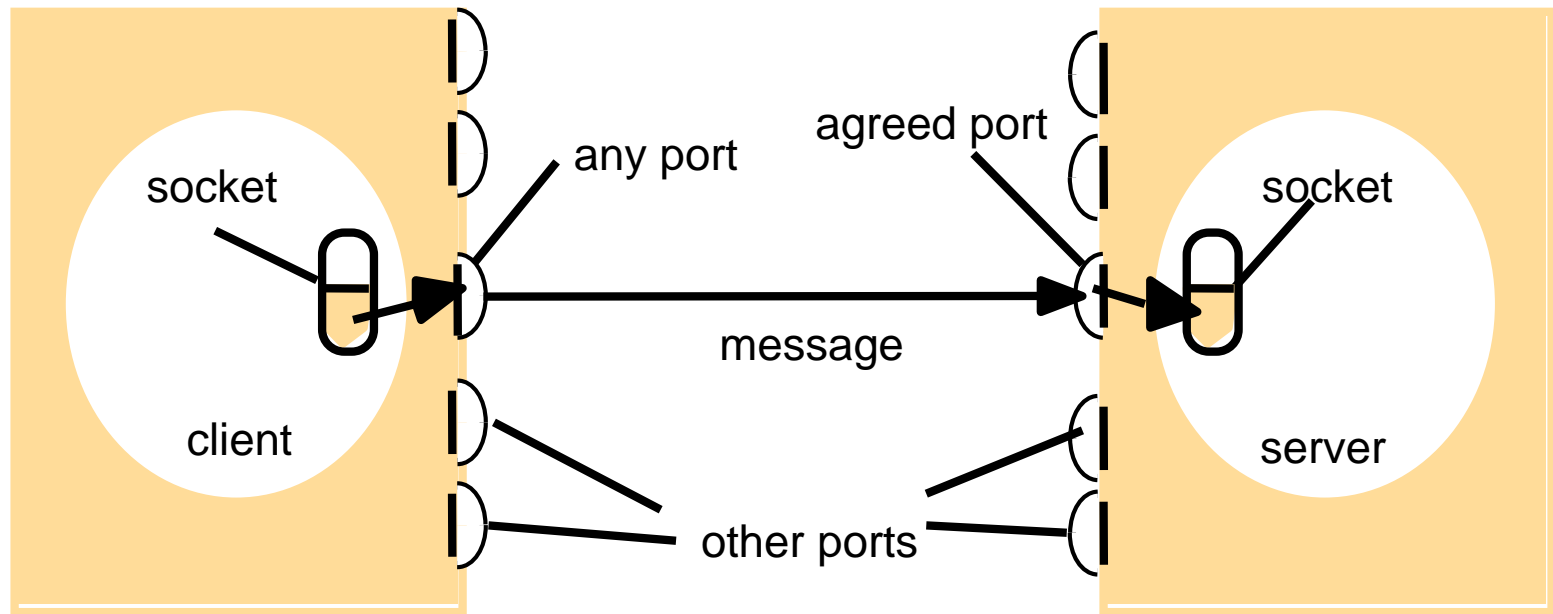


The characteristics of interprocess communication

- Messages are sent to (**Internet address, local port**) pairs.
 - A port has exactly one receiver but can have many senders.
 - Processes may use multiple ports to receive messages.
 - Any process that knows the number of a port can send a message to it.
 - Each computer has a large number (2^{16}) of possible port numbers for use by local processes for receiving messages.
 - Process cannot share ports with other processes on the same computer.
- **Validity**: A point-to-point message service is reliable if messages are guaranteed to be delivered despite packet drop
- **Integrity**: Messages must arrive uncorrupted and without duplication
- **Ordering**: messages be delivered in sender order

Sockets

- Interprocess communication consists of transmitting a message between a socket in one process and a socket in another process.
- Processes may use the same socket for sending and receiving messages.



Internet address = 138.37.94.248

Internet address = 138.37.88.249



Java API for Internet addresses

- Java provides a class, `InetAddress`, that represents Internet addresses. Users of this class refer to computers by Domain Name System (DNS) hostnames.
- For example, to get an object representing the Internet address of the host whose DNS name is `bruno.dcs.qmul.ac.uk`, use:
 - `InetAddress aComputer = InetAddress.getByName("bruno.dcs.qmul.ac.uk");`
- This method can throw an `UnknownHostException`.
- The interface does not depend on the number of bytes needed to represent Internet addresses



UDP datagram communication

- Datagram sent by UDP is transmitted from a sending process to a receiving process **without acknowledgement or retries**.
- If a failure occurs, the message may not arrive.
- To send or receive messages a process must first create a socket bound to an **Internet address** of the local host and a **local port**.
- A client binds its socket to any free local port.
- Server will bind its socket to a known **server port**.
- The receive method returns the Internet address and port of the sender, in addition to the message, allowing the recipient to send a reply.



UDP datagram communication

- **Message size:** The receiving process needs to specify an array of bytes of a particular size in which to receive a message.
- **Blocking:** Sockets normally provide non-blocking sends and blocking receives for datagram communication
- The send operation returns when it has handed the message to the underlying UDP and IP protocols.
- **Queue:** On arrival, the message is placed in a queue for the socket that is bound to the destination port.
- The message can be collected from the queue by an outstanding or future invocation of receive on that socket.
- Messages are discarded at the destination if no process already has a socket bound to the destination port.



UDP datagram communication

- **Timeouts:** Use to avoid waiting forever.
- UDP uses a **checksum** to ensure that the message is not corrupted.
- UDP datagrams suffer from the following failures:
 - **Omission failures**
 - **Ordering**
- A reliable delivery service may be constructed using acknowledgements which is not available in UDP.



Java Classes for UDP datagrams

- **DatagramPacket**: This class provides an instance that contains the message, length, IP and port
 - The message can be retrieved with the method `getData`.
- **DatagramSocket**: This class supports sockets for sending and receiving UDP datagrams. Methods include:
 - The constructor accepts a port. An exception(`SocketException`) is thrown if the chosen port is already in use or is reserved
 - `send` and `receive`
 - `setSoTimeout`



UDP client sends a message to the server and gets a reply

```
import java.net.*;
import java.io.*;
public class UDPClient{
    public static void main(String args[]){
        // args give message contents and server hostname
        DatagramSocket aSocket = null;
        try {
            aSocket = new DatagramSocket();
            byte [] m = args[0].getBytes();
            InetAddress aHost = InetAddress.getByName(args[1]);
            int serverPort = 6789;
            DatagramPacket request = new DatagramPacket(m, m.length(), aHost, serverPort);
            aSocket.send(request);
            byte[] buffer = new byte[1000];
            DatagramPacket reply = new DatagramPacket(buffer, buffer.length);
            aSocket.receive(reply);
            System.out.println("Reply: " + new String(reply.getData()));
        }catch (SocketException e){System.out.println("Socket: " + e.getMessage());}
        }catch (IOException e){System.out.println("IO: " + e.getMessage());}
    }finally {if(aSocket != null) aSocket.close();}
}
```



UDP server repeatedly receives a request and sends it back to the client

```
import java.net.*;
import java.io.*;
public class UDPServer{
    public static void main(String args[]){
        DatagramSocket aSocket = null;
        try{
            aSocket = new DatagramSocket(6789);
            byte[] buffer = new byte[1000];
            while(true){
                DatagramPacket request = new DatagramPacket(buffer, buffer.length);
                aSocket.receive(request);
                DatagramPacket reply = new DatagramPacket(request.getData(),
                    request.getLength(), request.getAddress(), request.getPort());
                aSocket.send(reply);
            }
        }catch (SocketException e){System.out.println("Socket: " + e.getMessage());}
        }catch (IOException e) {System.out.println("IO: " + e.getMessage());}
    }finally {if(aSocket != null) aSocket.close();}
}
```



Why Would Anyone Use UDP?

- Finer control over what data is sent and when
 - As soon as an application process writes into the socket
 - ... UDP will package the data and send the packet
- No delay for connection establishment
 - UDP just blasts away without any formal preliminaries
 - ... which avoids introducing any unnecessary delays
- No connection state
 - No allocation of buffers, parameters, sequence #s, etc.
 - ... making it easier to handle many active clients at once
- Small packet header overhead
 - UDP header is only eight-bytes long

Popular Applications That Use UDP

- Multimedia streaming
 - Retransmitting lost/corrupted packets is not worthwhile
 - By the time the packet is retransmitted, it's too late
 - E.g., telephone calls, video conferencing, gaming
- Simple query protocols like Domain Name System
 - Overhead of connection establishment is overkill
 - Easier to have application retransmit if needed

